

## Assingment 5

```
In [1]: import pandas as pd
import numpy as np
from sklearn import metrics
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
In [2]: from sklearn.datasets import load_boston
boston = load_boston()
```

```
In [3]: data = pd.DataFrame(boston.data)
data.head()
data.columns = boston.feature_names
data.head()
```

Out[3]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33

```
In [4]: data['MEDV'] = boston.target
```

```
In [5]: data.shape
```

Out[5]: (506, 14)

```
In [6]: data.columns
```

Out[6]: Index(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX', 'PTRATIO', 'B', 'LSTAT', 'MEDV'], dtype='object')

In [7]: data.dtypes

```
Out[7]: CRIM      float64
        ZN       float64
        INDUS    float64
        CHAS     float64
        NOX      float64
        RM       float64
        AGE      float64
        DIS      float64
        RAD      float64
        TAX      float64
        PTRATIO  float64
        B        float64
        LSTAT    float64
        MEDV     float64
        dtype: object
```

In [8]: data.isnull().sum

```
Out[8]: <bound method NDFrame._add_numeric_operations.<locals>.sum of          CRIM      ZN
INDUS  CHAS    NOX    RM    AGE    DIS    RAD    TAX  \
0   False False False False False False False False False False
1   False False False False False False False False False False
2   False False False False False False False False False False
3   False False False False False False False False False False
4   False False False False False False False False False False
..    ...    ...    ...    ...    ...    ...    ...    ...    ...
501  False False False False False False False False False False
502  False False False False False False False False False False
503  False False False False False False False False False False
504  False False False False False False False False False False
505  False False False False False False False False False False

        PTRATIO      B  LSTAT  MEDV
0      False False False False
1      False False False False
2      False False False False
3      False False False False
4      False False False False
..      ...    ...    ...    ...
501  False False False False
502  False False False False
503  False False False False
504  False False False False
505  False False False False

[506 rows x 14 columns]>
```

In [9]: data[data.isnull().any(axis = 1)]

```
Out[9]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	MEDV
--	------	----	-------	------	-----	----	-----	-----	-----	-----	---------	---	-------	------

In [10]:


data.describe()

Out[10]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.00
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	68.574901	3.79
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	28.148861	2.10
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000	1.12
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.025000	2.10
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	77.500000	3.20
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	94.075000	5.18
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000	12.12

```
In [11]: corr = data.corr()
corr.shape
plt.figure(figsize=(20, 20))
sns.heatmap(corr, cbar = True, square = True, fmt = '.1f', annot = True, annot_kv
```

Out[11]: <AxesSubplot:>



```
In [12]: x = data.drop(['MEDV'], axis = 1)
         y = data['MEDV']
```

```
In [13]: from sklearn.model_selection import train_test_split
```

```
In [14]: X_train, X_test, y_train, y_test = train_test_split(x, y, test_size = 0.3, random
```

```
In [15]: from sklearn.linear_model import LinearRegression
```

```
In [16]: lm = LinearRegression()
```

```
In [17]: lm.fit(X_train, y_train)
```

```
Out[17]: LinearRegression()
```

```
In [18]: lm.intercept_
```

```
Out[18]: 36.35704137659479
```

```
In [19]: coefficients= pd.DataFrame([X_train.columns, lm.coef_]).T
coefficients = coefficients.rename(columns = {0: 'Attribute', 1: 'Coefficients'})
coefficients
```

Out[19]:


	Attribute	Coefficients
0	CRIM	-0.12257
1	ZN	0.055678
2	INDUS	-0.008834
3	CHAS	4.693448
4	NOX	-14.435783
5	RM	3.28008
6	AGE	-0.003448
7	DIS	-1.552144
8	RAD	0.32625
9	TAX	-0.014067
10	PTRATIO	-0.803275
11	B	0.009354
12	LSTAT	-0.523478

```
In [20]: y_pred = lm.predict(X_train)
```


```
In [21]: print('R^2:',metrics.r2_score(y_train, y_pred))
print('Adjusted R^2:',1 - (1-metrics.r2_score(y_train, y_pred))*(len(y_train)-1)/
(len(y_train)-X_train.shape[1]-1))
print('MAE:',metrics.mean_absolute_error(y_train, y_pred))
print('MSE:',metrics.mean_squared_error(y_train, y_pred))
print('RMSE:',np.sqrt(metrics.mean_squared_error(y_train, y_pred)))
```

```
R^2: 0.7465991966746854
Adjusted R^2: 0.736910342429894
MAE: 3.089861094971131
MSE: 19.073688703469028
RMSE: 4.367343437774161
```

```
In [22]: plt.scatter(y_train, y_pred)
plt.xlabel("Prices")
plt.ylabel("Predicted prices")
plt.title("Prices vs Predicted prices")
plt.show()
```




```
In [23]: plt.scatter(y_pred,y_train-y_pred)
plt.title("Predicted vs residuals")
plt.xlabel("Predicted")
plt.ylabel("Residuals")
plt.show()
```



```
In [24]: sns.distplot(y_train-y_pred)
plt.title("Histogram of Residuals")
plt.xlabel("Residuals")
plt.ylabel("Frequency")
plt.show()
```

C:\Users\DELL\AppData\Local\Programs\Python\Python39\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)





```
In [25]: y_test_pred = lm.predict(X_test)

acc_linreg = metrics.r2_score(y_test, y_test_pred)
print('R^2:', acc_linreg)
print('Adjusted R^2:', 1 - (1 - metrics.r2_score(y_test, y_test_pred)) * (len(y_test) - 1) / (len(y_test) - 2))
print('MAE:', metrics.mean_absolute_error(y_test, y_test_pred))
print('MSE:', metrics.mean_squared_error(y_test, y_test_pred))
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, y_test_pred)))
```

R^2: 0.7121818377409185

Adjusted R^2: 0.6850685326005702

MAE: 3.8590055923707407

MSE: 30.05399330712424

RMSE: 5.482152251362985